

GPGPU Programming with CUDA

Introduction for dummies from dummy

Leo Kruglikov

Personal notes

2022

Contents

1	Basics of Architecture	5
1.1	Execution abstraction	5
1.2	Parallel execution and warps	7
1.3	Memory model	7
1.4	Memory allocation model	10
2	Programming in CUDA	11
2.1	Setup	11
2.1.1	Hello from CUDA	11
2.2	Threads & blocks indexing	12
2.3	Memory	13
2.3.1	Vector addition	13
2.3.2	Multiple dimension memory	16
2.4	More on Memory	18
2.4.1	Shared memory vs global memory	18
3	Basic parallel algorithms & patterns	21
3.1	Reduce	21
4	CUDA synchronization mechanisms	28
5	Streams	33
5.1	Concept of streams	33
5.2	Basic usage of streams	34
5.2.1	Asynchronous streams	34
6	Appendix	37

Author's preword

These *notes* are a kind of a collection of different articles from diverse resources on this topic. More precisely, the authors interpretation of these. A big part of the code snippets are also taken from different resources, and has not always been tested. The author will do its best to try to cite the sources. Therefore it is really a *collage* of notes, articles, books on the CUDA programming. The author's main goal is to provide the most detailed possible explanation of various code snippets, as well as try to explain the main features of CUDA programming.

Note that this document was initially written for the author itself, who is a physics major and is a fully self-taught guy in programming. For the author, it was a way of learning the topic and memorize the important concepts of it.

The goal of these notes is to **give us a good basic understanding of the GPU architecture, and the most importantly, try to fully depict the most common examples of CUDA codes.**

Dictionary

- GPU - Graphics Processing Unit
- CUDA - Compute Unified Device Architecture. The language we use to *talk to the GPU*. I will often refer to it as the CUDA API. In fact, it is not a language, but an API.
- Device - the GPU, from the software viewpoint. You may think of the notion of the device as an external executor of a function, in our case, the GPU.
- Host - the CPU, from the software viewpoint. The *machine*, that will launch GPU code from a usual C/C++ (or any other language) program, which, by default, would have been executed on the CPU.
- Kernel - nothing more than a function, that will run on the device(GPU).
- SIMT/SIMD - Single Instruction Multiple Threads/Data.

Small introduction

If one wants to perform computations on the GPU, one must have a way to address it. There are various API's developed. The biggest ones are the Khronos Group's OpenCL, Microsoft's Direct Compute, and the one discussed here, the Nvidia's Compute Unified Device Architecture, or shortly - CUDA. Do not mix it up with OpenGL, which is a slightly different thing, as it operates more on the graphics functionality.

When discussing the necessity of the GPU for computations, many come up with the example of the car and the bus [2]. Suppose you need to transport people from a point A to a point B . To solve this problem, you are given a car and a bus. What would be the most optimized way to transport these people? We introduce here the notion of throughput (bandwidth) and latency. The ability to perform a certain number of operations in a certain period of time is the throughput, and the amount of time that is required to perform a single operation is the latency. In our analogy, the bus, having a smaller speed than the car, but a greater capacity, has a big latency but a big throughput. On the other hand, the car has a small latency and small throughput.

So going back to our problem, we have that if the number of people to transport is significant, then the wise way to transport them is to use the bus. However, if the number of people is small enough, one should use the car, to get the small group of people faster to the point B . In this analogy, the car is the CPU, and the bus is the GPU.

I am convinced that after some examples of code using CUDA, the reader will understand, how powerful actually the GPGPU model is for certain tasks accomplishment. Like in our example with the bus and the car.




Figure 1: Schematic difference in architecture between the CPU and the GPU. Without going into details (~~as mentionned in the disclaimer, the author have not studied it in full depth~~), one is able to see that the GPU has many smaller ALU's. They are less powerful than those of the CPU and don't stand a chance in a theoretical 1v1 battle, but may do enough damage, when working together. [8]

1 Basics of Architecture

Before starting to consider some C/C++ CUDA code examples, we will look into some architecture of the GPU. Indeed, one of the differences between CUDA (or GPU) and usual/sequential programming (in our understanding, the *usual* programming is the code we write in C, Java, Python, etc...), is that one must take into account the architecture of the GPU, even when writing some simple code. The GPU has a multithread architecture by default, so when the programmer is partitioning the parallel tasks, he must make sure that there is no any redundant operations, and think about the way the cores will execute these tasks. If this partitioning takes into consideration all necessary aspects of the architecture and memory, it is possible to archive significant performance improvments.

The main difference between the CPU and the GPU is that the GPU has, in a way, lots of smaller CPU's in it, which are much less powerful than the actual CPU (Figure 1). [8]

1.1 Execution abstraction

As you might have noticed, the GPU is by definition a multi-threaded device. This means, it is suitable for the so called SIMT or SIMD (remember the bus and car analogy).

From the hardware viewpoint, we are distinguishing the **Device (GPU)** itself, the **Streaming Multiprocessors (SM's)**, and the **CUDA cores**. These are physical entities, having a certains structure and characteristics. The goal is not to give a detailed description of the GPU architecture, but rather to provide the idea of the CUDA mapping between the hardware and software world. While launching a kernel on the device, every mentionned part will be assigned a certain role, and will treat the software abstractions accordingly.

For us, programmers, we are writing software, and are operating with software abstractions. However, we still need to know how are these abstractions mapped onto the *hardware world*.

Threads are fundamental units of any GPU program. It is the most primitive *executor* of a function launched on the GPU. Threads (from the software side) are executed on the **CUDA cores** (the hardware side of the program).

Blocks are grouping entities that enclose threads. When a function is asked to run on the GPU, the blocks, which *contain* threads are delegated to the corresponding **Streaming Multiprocessor** or **SM**. So by now, we get that

$$\begin{array}{ccc} \text{Block of threads} & \xrightarrow{\text{are transmitted to}} & \text{SM} \\ \text{Threads in the block} & \xrightarrow{\text{are executed on}} & \text{Cores} \end{array}$$

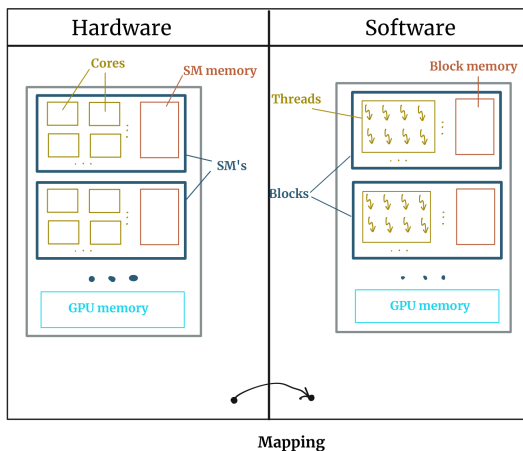
So we get that the SM's are partitioning the execution of threads on the Cores at runtime. For example, suppose we have launched 8 blocks of lets say 32 threads each. Suppose our GPU has 2 SM's. Then, as mentionned above, the blocks are divided and delegated to SM's. Thus for a GPU with 2 SM's, each SM will contain $8\text{blocks}/2\text{SM's} = 4\text{blocks}$, but if our GPU has 4 SM's, each SM will contain $8\text{blocks}/4\text{threads} = 2\text{blocks}$.

Note that from the programmer's viewpoint, the threads are strictly partitioned in blocks. From the hardware's viewpoint, however, it is not said, that threads *are physically divided into blocks*. This is the exact reason, we are discussing about *two worlds* and their mappngs.

Grid is the top-level abstraction layer from the software's perspective. The grid is the grouping entity that encapsulates blocks. We are thus considering that we are launching the grid on the **device**.

$$(\text{Device} \xrightarrow{\text{contains}}) \text{Grid} \xrightarrow{\text{contains}} \text{Blocks} \xrightarrow{\text{constains}} \text{Threads}$$

The mapping abstraction So to recap the mentionned notions, consider the sketch of the hardware/software mapping.



The abstraction between the hardware and the software side of CUDA is shown Figure 2. Once the function is provided, the programmer should think of the execution pipeline through threads, blocks, and the grid.

Figure 2: Hardware-software abstraction

1.2 Parallel execution and warps

We briefly saw the anatomy and the terminology of some underlying elements of the CUDA kernel execution. Conceptually, the threads, to whom a kernel was assigned execute in parallel and are grouped into thread blocks. Thread blocks run concurrently with each other grouped into a grid. It is important to note (see section 1.1) that the SM's will *automatically* assign the block's execution based on the GPU resources. One may say that *there is no promise on the block's concurrent execution*. [8]. So we do not know the order in which the blocks will be run.

However, there is a notion, which more or less guarantees the order of the execution of threads. **The Streaming Multiprocessor treats threads in groups of 32, which are called warps**. Think of the warps as a way to handle the threads, rather than a way of grouping (as the blocks of threads) ¹. We will see that the warps are an extremely important concept of the GPU development. We will see that the execution of kernel by threads is more efficient, when we take into account the fact that they are grouped by warps.

1.3 Memory model

The memory model of the GPU is quite complicated. It has different fields of memory that have different characteristics- latency of access, write/read modes, size, scope (to whom it is visible), etc... First let's take a look into **general** notions concerning the memory model.

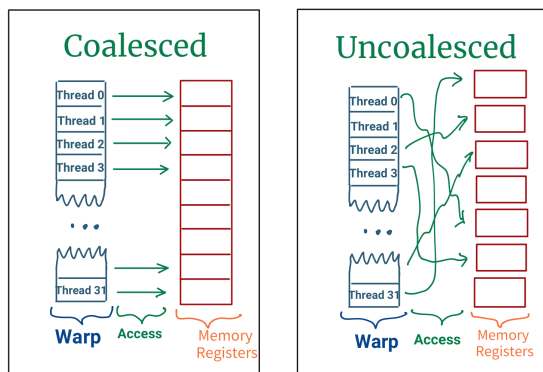


Figure 3: Memory access optimization mechanism.

beginning, but let's see how Nvidia is describing it [1]:

Global memory instructions support reading or writing words ² of size equal to 1, 2, 4, 8, or 16 bytes. Any access (via a variable or a pointer) to data residing in

¹In the AMD terminology, a warp is referred to as the wavefront. It brings more insight into the nature of warps

²Words can be data type of a certain size.

Coalesced vs uncoalesced memory access. Imagine a certain number of warps are scheduled by the SM. Lets say 2 blocks of 128 threads, which gives $2 \text{ blocks} \cdot 128 \text{ thr.} / 32 = 8 \text{ warps}$. These warps fetch some data from a certain place in memory of the GPU. We know that the warp is something very grouped. Therefore it would be nice that they access adjacent memory addresses. This notion may seem quite confusing in the

*global memory compiles to a single global memory instruction **if and only if** the size of the data type is 1, 2, 4, 8, or 16 bytes and the data is naturally aligned (i.e., its address is a multiple of that size). If this size and alignment requirement is not fulfilled, the access compiles to multiple instructions with interleaved access patterns that prevent these instructions from fully coalescing.*

Do not pay attention to the notion of **global** memory (we will discuss it soon). Try to read the Nvidia standard above again by looking at the coalesced scheme (Figure 3) to fully understand the mechanism. One may notice that this notion is one of the most crucial in the performance of the code. Indeed, when writing the kernel, one must keep in mind this aspect and try to ensure (when possible) a coalesced memory access.

Global memory of the GPU is the largest memory in terms of the size, and yet with the greatest latency. As we've discussed, the kernels are launched **from** host (the CPU). It would be wise to be able to share resources between the host and the device. For example, send data to the GPU from a usual C program and retrieve back in a processed form, otherwise the GPGPU programming does not make sense. This is exactly the purpose of the global memory. The global memory, as the name suggests, is *global*, i.e. it is **visible to all threads from all blocks**. As we will see in practice, the usual workflow of the program is to copy the data from the global memory to some other (which is discussed below), which is faster³ to manipulate.

Shared memory is much faster than the global one, but evidently smaller. Other crucial difference between global, is the scope - the shared memory is **only seen by threads in the same block**. This provides the ability for threads of the same block to share results and temporary calculations, and process the data IN PLACE. Think of the following situation: a grocery store, where the customers are threads. Every time a person wants to cook something at their house, they don't drive to the store to buy every ingredient needed. They rather go there once a week, for example, and buy the amount they need. They also make sure, that everything can fit in the fridge. Thus in this wonderful analogy, the fridge is the low latency shared memory, and the grocery shop - the big and unwieldy global resource - global memory. One sometimes refer to this memory as *cache memory controlled by the programmer*. However, it is important to take note that the reduced latency of the shared memory does not guarantee a better performance. Indeed, the biggest pitfall for all of us beginners are the *bank conflicts*⁴.

³You may think of it as the `malloc()` or `calloc()` functions in C or the keyword `new` in C++. Indeed, the allocation and the access to those variables is slower than declaring on the stack:

```
int * ptr_a = new int; is slower than int a = {};
```

⁴A small disclaimer: the notion of *bank conflicts* was one of the reasons for these personal notes, as it took a very long time, for the author to understand this concept. The reader should not panic if he's missing something. The examples will be discussed later in the practice part. So one should, if necessary, come back to this "theoretical" part after going through the examples. ~~The author wants to apologize for the eventual wordiness.~~

Bank conflicts. We already discussed the notion of warps, as an execution entity encapsulating threads. One may think of the banks as the analogy of warps in memory (i.e. warps are located at the *execution level abstraction*, and the banks-at the *memory level abstraction*). Shared memory is organized into banks. One *layer bank* is a sequential field of 32 memory adresses of 4 bits ($32 \# \text{adreses} \cdot 4 \text{bits}$).

Memory can serve as many simultaneous address as it has banks.

This is a very important property, so lest's consider an another illustrative analogy. Suppose in a national institution, each employe is assigned a counter, such that a single employe can serve only one client. Suppose you are the host (the person, who assigns people to desks) and you have a large group of people, who's number is the exact number as the number of desks. The wise choice would be to partition them between all the counters, right? Wouldn't that be silly to partition ,let's say 3, to one counter, 4 to an another, etc... and there will be, let's say 8 free counters. At these 8 counters, the employes will simply wait for customers, while there are customers, who are waiting in the queue. The analogy ~~may~~ not be the best, but the sketch should do the trick:




Figure 4: Memory banks serving the threads. Only one thread can access a bank with a certain ID simultaneously. From the analogy, threads are clients, and banks are employes at countes, giving them memory.

Read-only memory* is, as its name suggests, can't be changed by the kernel's threads, but is loaded at compile time. This memory is not as commonly described in the GPGPU programming documentation. It is often encapsulated in the GPU APIs, such as OpenGL. Thus by using those API's, the programming is indirectly reffered as the read-only memory.

Local registers By looking at the scheme of the architecture of the GPU vs the CPU (Figure 1), one may notice the amount of registers in the GPU vs the CPU. The amount of registers in the GPU is incomparable with the registers of the CPU. As you might have guessed, a register is memory, with the scope of a thread. The compiler of a CUDA program will try to optimize the number and size of registers. It is nevertheless possible, that the amount of memory in registers may fall short. Then the L1 and/or L2 caches will enter the play. This is the fastest memory available in the CUDA API, yet the most restrictive.

1.4 Memory allocation model

By now, we've made the difference between kinds of memory in terms of the scope, i.e, **who** and **when** is able to access various kinds of memory. This, of course, somehow impacts the way we're allocating it. However, one can classify the memory, in terms of its **allocation** procedure. There are mainly four ways of memory allocation.

We will see further that a standard program, which uses GPU resources, follows a certain path/pipeline. Roughly that is:

1. Resource declaration on the host (using `malloc`) and initialization using `memcpy`.
2. Memory allocation and initialization on the device.
3. Memory copying/transfer from the host to device.
4. Accelerated computation on the GPU and storage of the results from the device
5. Copy of data from the device back to the host.

One of the steps is the allocation of memory on the GPU from the host (by calling a certain CUDA function, which will be described later).

- Pageable memory
- Pinned memory
- Mapped memory
- Unified memory

The difference in terms of the API calls, use cases and mechanisms, will be explained further [REEEFFFF].

2 Programming in CUDA

By now, we have talked about the architecture of the GPU ⁵ by briefly discussing different notions - the execution model of threads and kinds of memory. Now, we will try to look at examples of CUDA codes and programs. We will try to refer to all these prior concepts in order to get a more detailed understanding. One may need to have a look at the section above to link the theory with the practice discussed here. I will do my best to choose the most illustrative examples (of those available in different resources) to explain specific concepts, as well as introducing some tools. Also note that I use Linux. (section).

2.1 Setup

Before starting writing some code, we must install the necessary packages and libraries. Every computer and operating system have their own subtleties, so the best way to install necessary tools, one should check the documentation for the specific system ⁶. For Linux, the main packages are `CUDA` and `CUDA-toolkits`. If we've already used our computer for some time, we probably have Nvidia drivers installed.

When we run a C program, we need a compiler - arguably, the most common one ~~and the best one~~ is `gcc`. For C++, we invoke his *improved* version `g++`. For CUDA programs, however, we need the NVidia's CUDA compiler - `nvcc`. As we've discussed, the program consists of host and device code. The device code is just plain C/C++ code, so `nvcc` is also able to compile plain C/C++ code. Note that CUDA code have a `.cu` extension.

```
$nvcc -o main main.cu
$./main
```

Listing 1: Compiling with nvcc and launching a CUDA program on Linux

If both of the programs do not return an error, the code has been compiled and launched successfully.

2.1.1 Hello from CUDA

Let's create our first CUDA and C++ program, and discuss every step below.

```
1 #include <stdio.h> //for printf()
2 #define N_THREADS 4 //number of threads
3 #define N_BLOCKS 2 //number of block
4
5 __global__ //declaration specifier
```

⁵Mainly about the Nvidia's architecture, but it can be quite well generalized to other GPU's.

⁶See the installation methods and required components on <https://docs.nvidia.com/cuda/cuda-installation-guide-linux/index.html>.

```

6 void hi_from_gpu(){ //kernel
7     printf("Hi from GPU\n");
8 }
9 int main(){
10     printf("Hi from CPU\n");
11     hi_from_gpu<<<N_BLOCKS,N_THREADS>>>(); //invoking kernel
12     cudaDeviceSynchronize(); //synchronize CUDA program
13     return 0;
14 }

```

In the output of the following code, we should get first Hi from CPU, followed by 8 times Hi from GPU. Let's discuss some aspect of the above code. Line 5 contains `__global__` declaration specifier which tells the compiler that the following kernel (function) can be launched on the device. In `main()` function, on line 11, we call the kernel. This is the semantics to invoke a CUDA kernel. The *arguments* in the angle brackets `<<<,>>>` are the dimension of threads blocks respectively that the kernel will be run on⁷ & . The type these dimensions can be either unsigned integer, or of the `dim3` type, which we will mention in a moment. In this case, the GPU will call 2 blocks, with 4 threads each. This indeed results in $8 = 2 \cdot 4$ invocations of `printf()`. Finally, from the main function, we call the `cudaDeviceSynchronize()` method, waits for all blocks before the main function returns.

2.2 Threads & blocks indexing

Now let's dive in deeper, and access threads and block indexing by modifying our `hi_from_gpu()` function:

```

#define N_THREADS 4
#define N_BLOCKS 2
__global__
void hi_from_gpu(){
    printf("Hi from GPU\n, from thread id %d and block id %d",
        threadIdx.x, blockIdx.x);
}

```

The output after executing from main,

```

$./main
Hi from GPU, from thread id 0 and block id 0
Hi from GPU, from thread id 1 and block id 0
Hi from GPU, from thread id 2 and block id 0
Hi from GPU, from thread id 0 and block id 1
Hi from GPU, from thread id 1 and block id 1
Hi from GPU, from thread id 2 and block id 1

```

We thus discovered how to access block and thread id's. The variables `threadIdx` and `blockIdx` are of type `dim3`, which is a simple structure, containing 3 unsigned int's. By

⁷We will see further, some additional stuff can be passed into these brackets, such as the size of memory and streams - concepts that will be seen in further sections

accessing its `.x` member variable, we are referring to the 1D indexing. In general, threads are indexed using `dim3` type. Thus, for a thread, living inside a block, there exists `x`, `y` and `z` dimensions. And same for the blocks, living in the grid - `x`, `y` and `z` components. Think of threads, We can launch many threads in many blocks, but how many? This depends on the hardware we are using. To get these specifications, one can look up this information online or asking our computer.⁸.

2.3 Memory

2.3.1 Vector addition

The first *useful* example that is usually introduced in CUDA tutorials is the vector addition. This example perfectly illustrates the need of GPU parallel model. Indeed, the component of the resulting vector x_i does not depend on other components. So the formula for vector addition is given by $x_i = a_i + b_i$. In sequential execution, we would create a loop, iterate over all elements and do something like `x[i] = a[i] + b[i]`. In order to parallelize this workflow, we must initiate N threads, with N = number of components in the vector. In C++, which simply runs on the host, we create the threads using the standard `std::thread`. We know, however, that the CPU does not have many threads (probably from 4 to 12 in most cases). This is where the GPU with its multithreaded architecture comes in. The idea here is to launch N threads which are spread over many blocks.

To discover memory allocation with code examples, we will consider vector addition. And to do that, one must first cover thread & block indexing for a more general case.

A more complex indexing. We know that the number of threads in one block is limited. Let's check on Nvidia's official CUDA documentation [1] concerning threads indexing.

For convenience, `threadIdx` is a 3-component vector, so that threads can be identified using a one-dimensional, two-dimensional, or three-dimensional thread index, forming a one-dimensional, two-dimensional, or three-dimensional block of threads, called a thread block. This provides a natural way to invoke computation across the elements in a domain such as a vector, matrix, or volume. The index of a thread and its thread ID relate to each other in a straightforward way: For a one-dimensional block, they are the same; for a two-dimensional block of size (Dx, Dy) , the thread ID of a thread of index (x, y) is $(x + y \cdot Dx)$; for a three-dimensional block of size (Dx, Dy, Dz) , the thread ID of a thread of index (x, y, z) is $(x + y \cdot Dx + z \cdot Dx \cdot Dy)$.

⁸To do so, one must find where `cuda` is located. In my case, it is in `/opt/cuda/`. Once here, we seek for `/samples/1Utilities/deviceQuery/` (or something similar). From here, we execute `./deviceQuery`.

Before writing some code, let's try to visualize the Nvidia's quote.




Figure 5: Simple example to illustrate a plain 1D indexing (supposing a block contains 4 threads). We see how the expression `BLOCKIDX.X*BLOCKDIM.X+THREADIDX.X` is used (see below). One can easily extrapolate the indexing to 2D and 3D cases, as described in the Nvidia documentation above.

Let's now write code to add two vectors, using a 2D indexing ⁹:

Ok, there are multiple aspects to discuss...

First let's discuss the code. First on the device. Remember we discussed about the global GPU memory. The line 31 we're working with memory. The classical pipeline of any CUDA program execution, is usually the following:

1. Begin code on host.
2. Allocate memory on host.
3. Allocate memory on device.
4. Calculations performed on the device.
5. Host treats the data.
6. `main()` gets returned.

Allocate memory on host. To allocate memory on host, we proceed with the usual C `malloc()` function and naturally casting the returned pointer to `double*` (don't forget to free memory at the end by `free()` function). Then we call a simple function `init_host_vector()` that will just initialize the data to some values, let's say to simple 0, 1, 2, 3,

Allocate memory on device. Then we allocate **global memory** on the GPU, by calling the `cudaMalloc()` function. Note the arguments that this function takes - a **pointer to pointer** - (`<type> **`) and the size in bytes.

⁹This is the first and almost the only **full** example of a cuda program. That is, the author will mostly provide the key aspects of the newly introduced features, thus skipping the usual parts of defining things, of allocating memory, freeing it, etc.

```

1  #include "stdio.h"
2  #define N_THREADS 512
3  #define N_BLOCKS 64
4  void init_host_vector(double *a, double *b);
5  void check_result(double *res);
6
7  __global__
8  void add_vec(double *a, double *b, double *res){
9      //compute the index
10     int id = blockIdx.x*blockDim.x+threadIdx.x;
11     if(id < N_THREADS*N_BLOCKS){
12         res[id] = a[id] + b[id];
13     }
14 }
15
16 int main(){
17     const int size_in_bytes =N_THREADS*N_BLOCKS*sizeof(double);
18     //initialize the data on HOST
19     //malloc() (C) or new (C++)
20     double *hst_a = (double *)malloc(size_in_bytes);
21     double *hst_b = (double *)malloc(size_in_bytes);
22     double *hst_res = (double *)malloc(size_in_bytes);
23
24     init_host_vector(hst_a, hst_b);
25
26     //allocate memory on GPU
27     double* dv_a;    cudaMalloc(&dv_a, size_in_bytes);
28     double* dv_b;    cudaMalloc(&dv_b, size_in_bytes);
29     double* dv_res;  cudaMalloc(&dv_res, size_in_bytes);
30
31     cudaMemcpy(dv_a, hst_a, size_in_bytes, cudaMemcpyHostToDevice);
32     cudaMemcpy(dv_b, hst_b, size_in_bytes, cudaMemcpyHostToDevice);
33
34     add_vec<<<N_BLOCKS, N_THREADS>>>(dv_a, dv_b, dv_res);
35     cudaDeviceSynchronize();
36     cudaMemcpy( hst_res, dv_res, size_in_bytes, cudaMemcpyDeviceToHost );
37
38     check_result(hst_res);
39
40     cudaFree(dv_res);    free(hst_res);
41     cudaFree(dv_a);     free(hst_a);
42     cudaFree(dv_b);     free(hst_b);
43     return 0;
44 }

```

Listing 2: Basic vector addition, using sequential thread indexing. [8]

Data copying between the host and the device is done using the `cudaMemcpy()` function. Also pay attention to the parameters - `cudaMemcpy(void* destination, void* source, size_t size, enum cudaMemcpyKind kind)`, with `kind`¹⁰ specifies how to copy. In code, it is understandable that the destination is device and the source is host.

Calculation. The kernel computes the sum of a vector of size $512 \cdot 64 = 32768$. Therefore 64 blocks of 512 threads are launched, and 32768 threads independently execute the `add_vec()` function. In the kernel, a unique ID is assigned to the thread (line 10) (see the indexing *policy* Figure ??). The ID's span from 0 to 32768, without ever repeating itself and covering all the numbers in the interval. We also add a simple `if()` statement, to be sure, that the thread's ID does not exceed the size of the vector. Thus every thread does exactly one calculation¹¹. Remember that the memory, in which these processes are happening, is the global memory.

Terminating. After the kernel, we call the usual synchronization function, which will, as usual, wait for all the threads and blocks to finish the calculations. Then, a simple check function on host, simply to ensure, that the parallel vector addition was successful and that we've received the correct result. And finally, the `cudaFree()` method, which does exactly what it is expected - frees the allocated memory on the GPU, the exact same things, as the `free()` does in a usual C code.

Let's also try to link this piece of code to the section about warps. Remember, threads are scheduled by the SM into warps - groups of 32 threads (subsection 1.2). In this case, we're launching a total of $512/32 = 16$ warps for each block, and 64 independent blocks. While working with CUDA threads, it is advised to work with the number, which is multiple of 32 (the thing we've done in this example). Indeed, imagine, if we were to launch 513 threads on a block. Then the SM would schedule 17 warps. This number of warps is the same as if we would launch 544 threads in a block. Indeed, this will launch $16 \cdot 32 = 512$ threads (executed simultaneously in one warp) **and** one additional thread on a almost empty/unoccupied warp.

One could also be interested in the difference in the execution time between 2 different configurations - either $64\text{threads} \cdot 512\text{blocks}$ or $32\text{threads} \cdot 1024\text{blocks}$. The benchmarking is very important and useful in GPGPU programming. Different tools and techniques of benchmarking CUDA programs will be discussed in further sections.

2.3.2 Multiple dimension memory

By now, we've looked into the global one-dimensional memory allocation (and copying) on device. Suppose you would like to work with 2D allocated memory. It is clear that we could work without the ability to explicitly allocate 2D or 3D arrays. Indeed, if we want to work with, let's say, a 2D matrix, we could simply transform it into a 1D vector of size $\text{rows} \cdot \text{cols}$.


¹⁰Here, I described the parameters in details. However, further on, the author will not be that precise. The reference for this function was taken from http://horacio9573.no-ip.org/cuda/group__CUDART__MEMORY_g48efa06b81cc031b2aa6fdc2e9930741.html. For further function descriptions, always have a look at it.

¹¹Of course, without taking into account the calculation of the id.

However, there is a feature in the CUDA API, that lets us to allocate 2D, and even 3D memory [4]. The method that implements 2D memory allocation is `cudaMallocPitch` with the following signature:

```
cudaMallocPitch( void** devicePtr,
                 size_t* pitch,
                 size_t widthInBytes,
                 size_t height);
```

The signature of function. Some new terminology terminology given on the right.



This function is allocating **at least** $\text{width} \times \text{height}$ bytes array. As we have seen, the allocation on the device is followed by the data copy from host. There is also a special function to do so - `cudaMemcpy2D()`. Apart from allocating and initializing the device data, we should be able to access it. So let's consider a code snippet that does so:

```
int main(){
    float *A, *dA;
    size_t pitch;

    A = (float *)malloc(sizeof(float)*N*N); // allocate on host
    cudaMallocPitch(&dA, &pitch, sizeof(float)*N, N); // allocate on device

    //copy memory
    cudaMemcpy2D(dA,pitch,A,sizeof(float)*N,sizeof(float)*N, N,\
        cudaMemcpyHostToDevice);
    /*...*/
}

__global__ void access_2d(float* devPtr, size_t pitch,\
    int width, int height) {
    for (int r = 0; r < height; ++r) {
        float* row = (float*)((char*)devPtr + r * pitch);
        for (int c = 0; c < width; ++c) {
            float element = row[c];
        }
    }
}
```

*Listing 3: To get this straight, one should know 2D array work in pure C. Indeed, to access an element, we first identify the current row, by doing some pointer manipulations. And then, once the pointer to the very beginning of the row is identified, we iterate over this row, by accessing it sequentially. Also note that in the code, in order to identify the row for the next iteration (i.e. it's 0's element pointer), one use the size of the pitch, and not of the width. We see that this 2D pitch is allocated **automatically** by the cuda memory management system. [3]*

As you might have guessed, this memory access is not extremely efficient as it is here, as here, really, all threads & blocks that the kernel will be launched with, will **all** access all the elements of the pitch memory. As one may say, THIS IS FOR EDUCATIONAL PURPOSES ONLY. In the documentation, it recommended, to use these functions to allocate and initialize the 2D memory on device. Hopefully, we all can find use cases of this memory and efficient indexing ¹².

2.4 More on Memory

2.4.1 Shared memory vs global memory

For now, in the examples, we've only looked into the global (& pinned) memory. It is global and accessible for all threads in all blocks for both read and write operations, and yet having high latency, comparing to more local memory types, such as the shared one. Shared memory's scope is one thread-block. Thus, if, let's say, a block contains 16 threads, the shared memory can only be read and modified by 16 local threads. If one wants to operate on it locally, the data must be copied to it. Therefore, if one wants this memory to be accessed/modified by 2 thread-blocks (= 32 threads), one must make sure that the data has been copied to the shared memory of the two blocks.

Let's now consider a more complex and yet classical example - matrices. Recall basic matrix operation -

- If A - $m \times n$ matrix (m rows and n columns), it can be added to some other matrix B of same size $m \times n$, by adding element-wise each element $a_{i,j} + b_{i,j}$.
- One can multiply two matrix A and B - $C = A \cdot B$ by using the formula $c_{i,j} = a_{i,1} \cdot b_{1,j} + a_{i,2} \cdot b_{2,j} + \dots + a_{i,n} \cdot b_{j,n}$, only if the matrix A is of size $m \times n$ and B of size $n \times p$. In short, the resulting element $c_{i,j} = \sum_{k=1}^{k=n} a_{i,k} \cdot b_{k,j}$ and yields C of size $m \times p$.

One can perform these operations using the global memory, or shared memory.

Global version code snippet is given down below.

Before moving to the code discussion, let's make a small *intermezzo* on transforming 2D array to 1D, to make things clear and easier for further analysis:

¹²There are a lot more of memory allocation, & memory copy methods, provided in the CUDA API. For example `cudaMallocArray()`, which has also a different signature. The different methods, dedicated to memory allocation and management can be observed here

```

1  #include "stdio.h"
2  #define BLOCK_SIZE 16
3
4  typedef struct{
5      int height; int width; float* element;
6  }Matrix;
7
8  __global__ void matmul_global(const Matrix a,\
9  const Matrix b, Matrix c);
10
11 int main(){
12     //init matrices A and B
13     Matrix A; A.height = 32; A.width = 32;
14     Matrix d_A; d_A.height = A.height; d_A.width = A.width;
15     int size = sizeof(float)*d_A.height*d_A.width;
16     cudaMalloc(&(d_A.element), size);
17     cudaMemcpy(d_A.element, A.element, size, cudaMemcpyHostToDevice);
18     /* same for d_B
19     ...
20     */
21
22     //prepare memory, for device to write to
23     Matrix d_C; d_C.height = d_A.height; d_C.width = d_A.width;
24     cudaMalloc(&(d_C.element), size);
25
26     //prepare dimensions of the kernel (2D indexing)
27     dim3 block_dim = (BLOCK_SIZE, BLOCK_SIZE); //dimension of block
28     dim3 grid_dim = (A.width/BLOCK_SIZE, A.height/BLOCK_SIZE); //dim. of blocks grid
29     matmul_global<<<grid_dim, block_dim>>>(d_A, d_B, d_C);
30
31     /*
32     cudaMemcpy(...); free(...); cudaFree(...); //free the resources
33     */
34 }
35
36 __global__ void matmul_global(const Matrix A, const Matrix B, Matrix C){
37     int row_id = blockDim.y*blockIdx.y + threadIdx.y;
38     int col_id = blockDim.x*blockIdx.x + threadIdx.x;
39
40     //accumulate sum for c_{row_id,col_id} element
41     float tempsum = 0.0;
42     for(int k = 0; k<A.width; k++){
43         tempsum += A.element[row_id*A.width + k]*\
44             B.element[k*B.width + col_id];
45     }
46     C.element[row_id*C.width + col_id] = tempsum;
47 }

```

Listing 4: Basic, yet important, global memory usage. [8]

Suppose we are given a 2D array (matrix). But we know that, behind the scenes, (even if we're accessing `a[i][j]` in many programming languages), it all breaks down to contiguous, linear memory addresses. So suppose you have allocated a 1D array `arr` of size N . Then, you could access your i 'th element by doing `arr[i]` (supposing C or C++) or by doing `*(arr + i)`, where `arr` - the address of the 0'th element of the array. So the formula is given by $A_{i^{th}} = A_0 + \text{sizeof}(\text{type}) \cdot i$, where A_0 - the first (0'th) memory address and sizeof - the size of one memory field (e.g. $\text{sizeof}(\text{char}) = 1$). Suppose now, that we've allocated a 2D memory array of size $M_{\text{rows}} \times N_{\text{cols}}$. For the first row one can apply our previous formula - $A_{0,j} = A_0 + \text{SZ} \cdot j$. However, if j is greater than N - the number of columns, a problem occurs. If we work in 2D indexing, over a loop, we would just *reset* our index j to 0 and increment i (if $j = N - 1$). However, in a flattened 2D case, in order to access $A_{i,j}$ address-wise, we use the expression $A_{i,j} = A_0 + \text{sizeof}(\text{type}) \cdot N \cdot i + j$. One can easily check, that this expression is consistent with all the examples. This expression is for row major arrays - accessing successive elements in a certain row.

Okay, now we can attack the code. From line 13 to 21, we are initializing the data and allocating memory. The lines 27 & 28 are initializing the dimensions (number of threads in x,y directions and the number of blocks in the x, y directions), that the kernel will be launched on. Note that we are using 2D indexing for the first time. In this case, `grid_dim` is the dimension of the grid, in which blocks are contained. The `block_dim` is the block dimension - the number of threads in the block. In this case there are $16 \cdot 16 = 256$ threads and $\frac{32}{16} \cdot \frac{32}{16} = 2 \cdot 2 = 4$ blocks. So there are $256 \cdot 4 = 1024$ threads partitioned between 4 blocks. This is exactly what we need, because the resulting matrix C is exactly 32×32 , which gives us 1024 elements. Therefore, if everything goes well, each thread will perform the calculation for each element $c_{i,j}$ of the matrix C . In theory, we would like that each thread iterates over one line of the matrix A and one column of B . Remember the 1D to 2D mapping. is going through all the elements in the row i of the matrix A (line 43).

$$\begin{aligned} A_{i,h} &= A_0 + i \cdot N_{\text{cols}(A)} + h \\ B_{h,j} &= B_0 + h \cdot N_{\text{cols}(B)} + j \end{aligned}$$

These equations (at least the first one) are exactly predicted by 1D to 2D mapping and are represented on the lines 43 and 44 of the code snippet. Finally, once each thread runs over its own row and column, it is gathering the result by assigning this accumulated sum to the `C[i][j] = C[i*Width + j]`, where `i,j` - row id's and column id's respectively¹³




Figure 6: Reduce algorithm with different strategies.

3 Basic parallel algorithms & patterns

3.1 Reduce

Consider the situation, when we need to add all the array's elements together. In C++, one could use the STL algorithm `accumulate()` function. In a more primitive implementation, we would do a single for-loop and accumulate the results in one variable. The complexity of such an algorithm would then be about $O(n)$, where n - the length of the array. However, one could add some parallelism to this algorithm. Indeed, at first iteration, what we would do is to add the 0th element with the 1st, the 2nd with the 3rd, the 4th with the 5th, etc... Notice that all of these $N/2$ additions can be done in parallel (that is, each thread does one addition). The next *iteration*, would be summing the result of the sum of 0th and 1st with the sum of 2nd and 3rd, thus giving a total of $N/4$ additions (instead of $N/2$). Therefore, the number of necessary divisions on the next step, will be the half of those, during the previous ones. As those divisions are performed in parallel, this algorithm looks like a log-scale complexity.

Global memory reduce

Let's first have a look at a not so naive implementation of the discussed reduce algorithm on the GPU. Once again, we're looking at a simplified version of the code, without implementing memory allocation, copy, etc... (note that in this case, we use simple global memory with `cudaMalloc()`). The implementation of this is given by 5.

Okay, let's discuss the 5. First in the host function, we define the number of blocks. In this case, this number is not so important. It would be important to optimize the execution, by taking into account the notion of warps, etc... The important part is the loop in the host code and the device kernel. We will try to do the debugger's job and inspect the steps. During the first iteration, the value of `stride` is 1. The important point is that the thread id, computed in the kernel does not depend on the value of the stride. This is because we're working with the global memory, and the access is global.

```

__global__ void reduce_global_kernel(float *data_out,\
                                   float *data_in, int stride, int size) {
    int idx_x = blockIdx.x * blockDim.x + threadIdx.x;
    if(idx_x + stride < size){
        data_out[idx_x] += data_in[idx_x + stride];
    }
}

void reduce_global(float *d_out, float *d_in, int n_threads, int size) {
    int n_blocks = (size + n_threads - 1) / n_threads;
    for (int stride = 1; stride < size; stride *= 2){
        reduce_global_kernel<<<n_blocks, n_threads>>>(d_out, d_in, stride, size);
    }
}

//main()

```

Listing 5: Global memory reduction. [8]

Suppose the size is 32, partitioned into 1 block. Then for the first iteration, we'll get, $out[0] = [0] + [1]$, $out[1] = [1] + [2]$, ... $out[30] = [30] + [31]$. This is exactly the first iteration, showed in Figure 7 (the case of $size=8$). Then, during the next iteration, we're *jumping* over 2 next elements, and adding them, in order to get the sum of N_{stride} elements, and save them into the $data_out[0]$. Let me mention, that this process is illustrated in the image above Figure 7.

This algorithm is, maybe, not easy to understand, but is very fundamental parallel algorithm. Both the algorithm, and the way of analyzing the problem. However, it can be optimized using the block's shared memory.




Figure 7: The description of every iteration, for the global memory reduction kernel. Note the how stride is doubling every iteration, and how the elements are accumulating in the very first (0th) element of the output data. One should understand, that the code works for arbitrary number of blocks, as we're working with global memory, visible to all threads in all blocks

Shared memory reduce

As we've discussed several times above, the shared memory access has low latency. The idea is thus to do a local copy of the data to each block. As the shared memory's size is limited, we **map** the data pieces to each block (see figure).

Now, let's try to understand it together in more detail, using well-defined numbers of threads and blocks, to make things more illustrative. Trying to keep in mind both the illustration (Figure 8).

```

1  /*Perform the necessary declarations, main(), before/after, etc...*/
2  void reduction(float *d_out, float *d_in, int n_thr, int size){
3      cudaMemcpy(d_out, d_in, size * sizeof(float), cudaMemcpyDeviceToDevice);
4      while(size > 1){
5          int n_bl = (size + n_thr - 1)/n_thr;
6          reduce_shared<<<n_bl, n_thr, n_thr*sizeof(float), 0>>>(d_out, d_out, size);
7          size = n_bl;
8      }
9  }
10
11 __global__ void reduce_shared(float* d_out, float* d_in, unsigned int size){
12     int idx_x = blockIdx.x * blockDim.x + threadIdx.x;
13     extern __shared__ float s_data[];
14     s_data[threadIdx.x] = (idx_x < size) ? d_in[idx_x] : 0.f;
15     __syncthreads();
16
17     // do reduction
18     for (unsigned int stride = 1; stride < blockDim.x; stride *= 2) {
19         // thread synchronous reduction
20         if ( (idx_x % (stride * 2)) == 0 ){
21             s_data[threadIdx.x] += s_data[threadIdx.x + stride];
22         }
23         __syncthreads();
24     }
25     if(threadIdx.x == 0){
26         d_out[blockIdx.x] = s_data[0];
27     }
28 }

```

Listing 6: Optimized reduce, with shared memory. [8]

We omit the `main()` function and the host/device memory allocation/copy. Line 2 declares the function, which will be calling the kernel. Line 3 copies memory `cudaMemcpyDeviceToDevice`. This is done in order to make sure we're working with the same memory locally allocated on the device. This is illustrated with the bottom blue arrow.




Figure 8: Reduce algorithm using shared memory.

First iteration For the first iteration, we're associating the variable `int size` to be the literal number of elements to be reduced. In our case, it is 32. The next declared variable `int n_bl` is, as the name suggests, the size of the block. As we're working with the shared memory, each shared memory will be allocated for one specific block, as shown in figure. In our case, we've chosen *nice* numbers, so that $n_thr * n_bl = size = 32$, with `n_thr` - the number of threads per each block. On line 6, we're finally invoking the kernel with the initial 8×4 dimensions (the first 2 parameters in the angle brackets). The 3rd parameter in the angle brackets is the size of shared memory, that will be assigned to each block. We see this syntax for the first time. This is how we allocate the **dynamic shared memory**, which we will discuss a bit later, as well as the last parameter 0 (ignore that too for the moment). For now, this is just shared memory allocation, outside the kernel itself. So we've allocated $n_{threads} \times size_{float}$ - the exact amount of shared memory, that will be accessed by the 4 threads.

Let's move to the kernel body itself. First, on line 12, our usual procedure, we're assigning a personal ID to each thread. Line 13 declares shared memory with size of $n_{threads} \times size_{float}$, specified outside, at the kernel call. This is the dynamic allocation of shared memory with `extern` keyword (as said, we'll discuss it later). On Line 14, we're copying the data to shared memory. Remember, the size of shared memory is the size of threads in the block (line 6). Thus the operation `SH_DATA[LOCAL_THR_ID]` is performed. Line 6 operation is illustrated with the violet lines - the mapping between the global to shared memory ($[0 : 3]_{global} \mapsto [0 : 3]_{block=0}$, $[4 : 7]_{global} \mapsto [0 : 3]_{block=1}$, ...). After copying, we are making sure that **all** the threads are done copying with `__syncthreads()`. Without that, some problems can occur (e.g. if we're starting reducing in 0'th block **before** all the threads within this blocks are done copying its data).

On line 17, we're starting to perform reduction. Let's first try to ignore the `if` condition on line 20, in order to better understand, why is it, and should be here. The dummy variable `stride` varies from 1 to the size of the block - $1 \rightarrow 2 \rightarrow 4 \rightarrow 8 \rightarrow 16 \rightarrow \dots$ (in our case till 8). For example, for a local thread with local ID = 0, the sum of `SH_MEM[0]`, `SH_MEM[1]`, `SH_MEM[2]` will be accumulated within the loop. For thread with local ID = 1, the sum of `SH_MEM[1]`, `SH_MEM[2]`, `SH_MEM[3]` will be accumulated. For thread with local ID = 2,

the sum of `SH_MEM[2]`, `SH_MEM[3]`, `SH_MEM[4]`, will be accumulated, and so on. We have now several problems, e.g. the access to `SH_MEM[4]`, which is out of bounds, as the size of it is the same as number of threads (4 in our case).

Consider now the loop on the global scale/scope, as on line 20, we're checking the global thread ID. Then which threads will access the line 21, for the dummy variable `stride = 1`? These are threads with global ID 0, 2, 4, 6, 8, 10, 12, 14, 16, ..., which corresponds to local threads $\{0, 2\}_{block=0}$, $\{0, 2\}_{block=1}, \dots$ (see violet lines Figure 8). On next iteration of the dummy variable `stride = 2`, only threads with global ID's 0, 4, 8, 12, 16 will access the line 21, which corresponds to local threads $\{0\}_{block=0}$, $\{0\}_{block=1}, \dots$

Now, taking into account the two aspects, we can say, that for the first iteration of the dummy variable `stride=1`, the threads with local ID's 0, 2, will accumulate `SH_MEM[0]`, `SH_MEM[1]` into `SH_MEM[0]` and `SH_MEM[2]`, `SH_MEM[3]` into `SH_MEM[2]` respectively. For the next iteration of the dummy variable `stride=2`, only the thread with local ID's 0 will accumulate `SH_MEM[0]` and `SH_MEM[2]` into `SH_MEM[0]`. **But remember:** in the previous iteration, the sum of `SH_MEM[0]`, `SH_MEM[1]` was stored in `SH_MEM[0]` and `SH_MEM[2]`, `SH_MEM[3]` into `SH_MEM[2]`. Thus, during the last iteration, we've performed the reducing operation within the block and accumulated the sum into `SH_MEM[0]`. The line 25 will simply write the value of `SH_MEM[0]` to the output, global memory. The kernel is done.

Next iterations operate the same way as the first. The only thing that is changing is the number of blocks, that the GPU will operate with. This does neither change the workflow, nor even the size of shared memory within each block.

It is quite hard ~~even maybe very hard~~ to understand the pipeline of the method execution. We should re-read the text above again and again, and try to associate it with the Figure 8. To recap the reduce process with shared memory:

- Define the initial number of threads and blocks, such that they cover the whole array to be reduced (note that `#threads` - number of elements on which the reduction will be performed locally).
- At every iteration, launch the kernel with the `#blocks` and update the new size of the array, which contains the previously reduced elements.
- In the kernel :
 1. Assign global thread id
 2. Copy data to the shared memory, from the global memory.
 3. Perform the reduction in the shared memory locally. With the if condition, make sure that
 - (a) The memory access does not overflow
 - (b) The elements are not added more than once (only add 2 consecutive elements)

4. Copy the data at 0'th location (the location of the elements accumulated within the block) to the global memory. (in Figure 8, this corresponds to the blue grid below)
- Repeat the kernel, by adjusting the size of the global array, (accessed by kernel at the beginning), controlled by the # of blocks.
 - End when the #blocks has reached 1 - when we're left with 1 block, on which we must perform the reduction and store at the 0'th element.

Okay, let's now discuss various performance aspects:

Memory. Clearly the main difference between the two implementations is the usage of memory. The first implementation uses global memory. Every thread goes to the global memory to take data, which, of course, takes time. In the shared memory implementation, the algorithm spends time to initialize the shared memory, which takes time. However, further on, it has lower latency than the global one. In general, the performance of the shared memory implementation is better than the global memory. Nevertheless, it is almost always advised to implement benchmarking into the code and/or use some debugging/benchmarking tools.

Warps. One may analyze the code under the warp's viewpoint. Remember, the threads are scheduled on the SM, partitioned into groups of 32 - warps. Ideally, they all run in parallel and don't have any barriers. Suppose that some threads in the scheduled warp, have some conditions that stops them, so they finish earlier than those, who haven't entered the condition. This means that some threads are idle, plus this requires additional, potential rescheduling. This problem, which causes throughput inefficiency, is called **WARP DIVERGENCE**. Let's now quickly try to detect warp divergence in both codes.

In the shared memory implementation, we've got one potential `if()` condition. Which may cause warp divergence. Indeed, the greater the stride is, the more threads will fail the `if(id_x+ stride<size)` condition, thus being idle & waiting for other threads, who have entered the condition.

The similar issue is in the second code. Indeed, on line 20, we have a condition, `if()`. The warp divergence is pretty big ¹⁴, as at every loop, only some threads will *enter* the condition (see the code discussion above), while other will become idle.

To attack these issues, one may use various techniques, potentially discussed in further sections. Some of these methods may be very tricky, sometimes requiring built-in CUDA features (unknown for us at the moment), and sometime very *primitive* techniques (e.g. section 6)

¹⁴One could potentially deduce the mathematical formulation of warp divergence, and evaluate, where is the divergence more present. But for the moment, we will stick with the qualitative approach.

To be fully honest, there is almost never a way to completely get rid of warp divergence. However, it is possible to do small changes to reduce them. In this case, we will follow a strategy, that has changed a bit the access of the elements and gather/reduce them together. Remember, in the shared memory implementation of reduce, we were looking for elements, which are located next to each other. We want to modify the memory access of threads, such that the pairs of elements are not necessarily next to each other. To do that, we are dividing the block in 2 parts, and we're adding (reducing) the 0'th element of the first half and the 0'th element of the second half. We call the dimension of the half of the block - the stride. Thus we get that `SH_MEM[0] = SH_MEM[0+stride]`, `SH_MEM[threadId.x] = SH_MEM[threadId.x+stride]`. Note that this expression may cause bad memory access, if `threadId.x + stride` is greater than the size of the shared memory. To prevent that, we're adding an additional condition - `if(threadId.x < stride)`. And this process will be done at every iteration of the for loop in the kernel. Here we are doing nothing but a literal reduction - *Dividing the block in half, adding(or any arbitrary operation) the one-to-one elements. When done, "throw" away the right block and perform the same reduce on the newly created block.* From the warp divergence perspective, one may notice that there is still a condition, that will potentially lead to warp divergence. However, looking at this condition, we can make a statement about **when** will this divergence occur. As the stride vary from `blockDim.x` to 0 being every time divided by 2 (e.g. 64, 32, 16, 8, 4, 2, 1, 0). The `if` condition will be omitted **if and only if** the warp size is less than the stride. Thus, at iterations, when the stride $\geq size_{warp}$ no warp divergence will occur, as **all the threads will pass the condition and there won't be idle threads**. As the warp size is 32, one can choose the most optimal block dimensions. Supposedly, the bigger the bloc dimension is, the less iterations will cause warp divergence, the better it is.

Personally speaking, this code/approach is much easier to understand and visualize than the previous ones, and in addition a bit faster. However, I wanted to roughly take some course/book's paths, where the reduction is presented in this specific order.

```

1  __device__ int reduce_sum(
2      cooperative_groups::thread_group gr, \
3          int *temp, int val){
4      int lane = g.thread_rank();
5
6      for (int i = g.size() / 2; i > 0; i /= 2){
7          //map each element in the first "semi" block
8          //to it's corresponding element in the second one
9              temp[lane] = val;
10             g.sync(); // wait for all threads to store
11             if(lane<i) val += temp[lane + i];
12             g.sync(); // wait for all threads in to load
13         }
14         return val; //only thread 0 will return full sum
15     }

```

Listing 7: This method, is ~~almost~~ the same as the first, optimized version of the reduce algorithm, using the shared memory `temp`. Therefore, one must note that this `reduce_sum()` method must be called for the array `temp*`, located in the shared memory.

4 CUDA synchronization mechanisms

As we've discussed in the section on architecture, when writing kernels, we need always to think about the GPU's hardware, scheduling, etc... By now, with the basic examples we've considered only *basic* operations. Indeed, the only API function we've used in the kernel is the `__syncthreads()`. This is a simple syncing mechanism, that ensures that all the threads **within the block** are done before this barrier. This is block-level synchronization barrier.

4.1 Cooperative Groups

Cooperative groups is a relatively new feature to the CUDA API. As the name suggests it, this feature enables us to group threads, with the ability to perform common, collective operations (or simply collectives). We can also perform synchronization between the threads, belonging to the same cooperative groups. With these API features, one can simplify the code, thus avoiding common mistakes and make it more readable. For example, in the code snippet section `??`, we perform the exact same algorithm as in `??`, by using some utility of the CUDA API. In this case, the `cooperative_groups::thread_group` class (do not pay attention to how we created this object and/or how it is declared). The `thread_rank()` function gets the ID/rank of the thread within the thread group `g` (the same way as `threadIdx.x` within a block). Then we're calling the `sync()` function, which ensures that all the threads within the thread group will be done setting the `val`, and the second to ensure that all threads are done reducing. It is important to understand, that all the threads will return the `val`. However, only the thread 0 will accumulate all the `val`'s. [3]

```

auto tb = this_thread_block(); // gets the thread block in kernel
tb.sync() // same method as in cooperative_groups
cooperative_groups::synchronize(tb);
this_thread_block().synchronize();
cooperative_groups::synchronize(this_thread_block());

```

Thread blocks

We've already seen the notion of a thread block many times. This notion was always quite *abstract*. Indeed, while launching the kernel, we've always kept the notion of thread blocks in our mind, but never actually accessed it. However, in newly introduced features, we can "access" the thread block explicitly. Remember the legacy `__syncthreads()` function. Well, syncing the `this_thread_block()` does the same thing as the `__syncthreads()`. Thus, there are several ways/semantics to synchronize the threads. The following function calls are synonyms.

One can also mention other synonyms `dim3 threadIdx` \equiv `dim3 thread_index()` and `dim3 blockIdx` \equiv `dim3 group_index()`. Thus one can easily replace these built-in keywords with these new methods, without any noticeable performance issues.

Partitioning

For these cooperative groups, a partitioning feature is also available. For instance, if we've created a thread block, by invoking `auto tb = this_thread_block()`, one can divide it into more small parts, for instance, into groups of 16 threads. This is done using the `cooperative_groups::partition()`, method, which takes the subject itself (the one to be partitioned into groups) and the number of threads per group. For instance, `cooperative_groups::partition(tb, 16)` divides the thread block into groups of 16 threads (so if e.g. a block has a max of 64 threads, this function will create) 4 groups of 16 threads in each.

The object returned is a `thread_group`. By accessing this object, it is possible to get the thread's rank, **within the obtained thread_group** (for instance, if we divide the thread block into groups of 16 threads and, by passing this object to a device function, print the `thread_rank()`, method, we will see numbers varying from 0 to 15).

The utility of these features overall is that it is less easier to make errors. Indeed, the NVidia documentation states that the usage of these features significantly reduces the risk of deadlocks. The concept of deadlocks is probably well known to the reader. This is a typical situation, when we don't want different threads to access a critical section, and ask them to be synchronized before accessing them. Consider the two pieces of code:

```

__device__ int sum(int *x, int n)
{
    ...
    __syncthreads();
    ...
}
__global__ void parallel_kernel(float *x, int n)
{
    if (threadIdx.x < blockDim.x / 2){
        sum(x, count); // Half of the threads enter and
                        //the other half-not
    }
}

```

```

__device__ int sum(thread_block block, int *x, int n)
{
    ...
    block.sync();
    ...
}
__global__ void parallel_kernel(float *x, int n)
{
    sum(this_thread_block(), x, count); //OK
}

```

Listing 8: Synchronizing using the legacy mechanism vs using the cooperative groups. [3]

In a nutshell, we clearly see a deadlock in the first piece of code, as there are only threads with ID's less than the half of the block dimension, which will enter the if condition. Those threads will perform their piece of code independently, and then will wait for all the other threads in the block to be completed and synchronized. This is a big issue, as there are threads, which will never start the `sum()` kernel and will just sit up there, waiting for those, who have. So the two chunks of threads will just wait for each other.

This is why one may find an application for the previously discussed primitives. In the second piece of code, one call the method `sum()` with a thread block. However, we could have divided into groups, using either CUDA functionality discussed above, **and only synchronize that block**, which, we are sure, will by run by all threads in the block.

Warp synchronizations

While programming with CUDA, one never gets tired to tired to think about warps, and how to optimize their execution. I do agree that it is not an easy task, to think of it, while doing even some basic operations. The new NVidia architectures and new versions of the CUDA API, provide a simple way to navigate through these concepts.

We have discussed a lot the advantages of the shared memory (e.g. for the speed and efficiency of the reduce algorithm). However, the new utilities give us a more faster, or

even a more local way to perform some operations. Remember, shared memory is block-local memory. Remember also that every thread has some kind of register to store small intermediate values while performing a kernel, for example, when we used to store the local thread ID. The so called *warp level synchronization primitives* allow us to access a certain thread's local register from another thread, as long as they are in the same warp, without the usage of the shared memory. Again, there are many things to keep in mind, but if such a function is called, it is doing everything *atomically*, in the sense that it is a primitive operation, performed locally on the threads in the warp. We therefore introduce here the notion of the **lane** (in fact, we used it briefly above). A lane is the thread id within the warp.

A little disclaimer: There are various Warp-level primitive functions. We will note that many function's name are similar, and only differ by the postfix `_sync()`. For instance, `__shfl_xor()` and `__shfl_xor_sync()`. Indeed, the ones with the `_sync()` postfix is a improvement of the former. It is recommended to use the newer version instead. I will not go into great details between these differences. I will just mention that there are differences in parameters ¹⁵(see further examples).

The `__activemask()` primitive/function is actually not a synchronization mechanism, but more of a filtering mechanisms. This function returns the *indices* of the active threads in the warp, where it was referenced from¹⁶. So the result returned from the `__activemask()` is used to call other synchronization functions, to *give them the corresponding threads, that are active in the warp*. So, for example, one could call the `__syncwarp(MASK)`, thus asking to sync all the threads meeting the `__activemask()` condition. One can go to the NVidia's developer's guide [1] and find the following:

Returns a 32-bit integer mask of all currently active threads in the calling warp. The Nth bit is set if the Nth lane in the warp is active when `__activemask()` is called. Inactive threads are represented by 0 bits in the returned mask.

So let's have a quick look and example at what did NVidia provide us with ¹⁷

- `__shfl_sync()` or in previous CUDA versions `__shfl()` is a tool, to "broadcast", or "spread" a certain value from a certain thread (identified with its lane) to all others in the warp. For example, in a certain warp, all the threads have a variable `int b = //some random int`, unique for all the threads. And I want all these

¹⁵Frankly speaking, the author hasn't seen this additional parameter being used in a very extensive way. This *extra* parameter is often replaced with some kind of hardcoded value

¹⁶One can think of the mask as the python's numpy feature, while doing a filter : `arr[:]<1.0`, which will return an array of booleans, which will be used to access later the elements of interest

¹⁷There are many such primitives available. As usual, the best way to get information about such CUDA functionality, is to look it up in the NVidia's developer's guide [1].

thread's variable `b`, to be the same as the one in the thread 4 (its lane number or ID within the warp). The best way to do that is to use the provided function: `__shfl_sync(0xffffffff, b, 4)` (or `__shfl(b,4)`). The second parameter is the variable to be broadcasted and the third one is the lane number to take the value from (because, of course, all the threads have this local variable `b`, which is different for all of them). So we're replacing all the `b`'s with THE `b` of the thread 4. The first parameter is actually the mask/filter, that tells the processor (or core I should say) which threads will be involved in this operation. This can be used by passing the result of e.g. `__activemask()` function (there are multiple *filtering* functions) or by passing it the *default* value in hex notation, which corresponds to the maximum number that can be displayed in binary notation (all the 32 bits are 1, thus we're saying that all the threads are active).

- `__shfl_up_sync()` or in previous CUDA versions `__shfl_up()` is a function to shift the values of the warp by an offset. For instance, let's say that in the warp, I want the 4th thread to have the value from 0th thread, the 5th thread the value of the 1st, the 6th thread the value from the 2nd thread, etc ... Then we would want to use the `__shfl_up_sync()` function, with the same parameters as in the `__shfl_sync()` function described above.
- `__shfl_down_sync()` Is the same idea as the `__shfl_up_sync()`. The difference is that we would use it if we wanted e.g. the thread 29 to have the value 31. (see figure for better understanding).

These primitives come in various shapes and forms. It would take quite a time to discuss them all here. The idea for them all, however, follows quite well the patterns, we've discussed just above. It is important to understand that these operations are sort of atomic, because, as we've seen, these are warp-local primitives, which is the most fundamental part of the execution scheduling model.

Atomics

When learning multithreaded programming, one of the first notions, that one must understand is the notion of atomics or atomic operations¹⁸. To make it short, these operations are operations, which, we are sure, will be performed *in the smallest possible period* and nothing will be able to affect the execution of these operations. For example, in C++, there is a `std::atomic<bool>` template specialization.

In the CUDA API, multiple atomic operations are provided. Arithmetic atomic operations, such as `atomicAdd()`, `atomicSub()`, `atomicMin()`, ... , bitwise functions, such as `atomicAnd()`, `atomicXor()`, ... (for more: [1]). All these operations have the same (or almost

¹⁸Intuitively, this comes from the word *atom*. This is a greek word, meaning something like *uncuttable*, *undivisible*. In fact, this is almost exactly, what an atomic operation corresponds to [5].

same) signatures: `int <or any primitive type> <function_name>(int* old_address, int value)`
 So these functions take the old address as the first parameter (for example, the address of the value, we want to add to) and the value as the second. The function thus performs an arithmetic or logical operation atomically, and stores the result in the old address. It also returns the value.

The atomic operations are however very expensive, as the scheduler must perform sequentialize the memory access. It is thus very advisable to use these operations, only when needed. For example, let's say we've performed some kind of reduction on multiple blocks locally. We know that if we've used some shared of memory on this reduction, one must in addition add the block results, to get the full answer of our reduction. One way to do that is to force an atomic operation in all blocks, so they store it in the first address of the global memory, that we will later *consume* from the CPU. So if one suppose that we've accumulated the thread-local reduce into `sum` and the pointer to the initial data `g_out`, we would do something like

```
if(this_thread_block().thread_rank()==0){
    atomicAdd(&g_out[0], sum);
}
```

5 Streams

5.1 Concept of streams

CUDA streams is an another fundamental concept that we've used previously in an implicit manner. Indeed, let me provide an another example using C++: when a "Hello World" programm is written in C++, no one cares about the thread usage. We do not create a "main" thread, when printing a character to the standart output. In the context of the CUDA API, **any** programm is by definition multithreaded. Therefore the analogy between the *usual* C/C++/Python/... and a GPU programm does not make sense. The stream concept is more of a thread **over the GPU threads**. By now, all the programs we've considered were launched in one single CUDA stream. So one can launch a programm (e.g. a simple vector addition) in one stream, which will be performed in parallel. However one can add a second vector addition programm, which will be done in parallel, with the first vector addition.

A CUDA stream is thus a way to make two (or more) kernels run in parallel. (Remember that the kernels themselves are already run in parallel following the CUDA architecture). By now, we've used one implicit stream - the default stream.

5.2 Basic usage of streams

Let's try to explicitly create a stream and look at its initialization syntax:

```
cudaStream_t stream;
cudaStreamCreate(&stream);
foo_kernel<<< grid_size, block_size, 0, stream>>>();
cudaStreamDestroy(stream);
```

Therefore to create the CUDA stream, one must first initialize the stream, then create the stream, i.e. allocate the memory via the `cudaStreamCreate()` method by passing the pointer of the stream. In order to call the kernel and make it run in a particular stream, one pass it as the last parameter in the angle brackets. If nothing is passed in the in the angle brackets for the stream, the default stream is used.

The simplest usage example of cuda streams is when one would want to concecutively call a certain kernel multiple times (for example N times). the way to do it using the default stream is to simply loop N times and call the kernel at every iteration. The way to do it using streams, is to create an array of streams and call the kernel by assigning it to different streams. The kernels will then be run in parallel without waiting the previous kernel return. The execution of the kernels So these functions take the old adress as the first parameter (for example, the address of the value, we want to add to) and the value as the second. The function thus performs an arithmetic or logical operation atomically, and stores the result in the old adress. It also returns the value.

The atomic operations are however very expensive, as the scheduler must perform sequentialize the memory access. It is thus very advisable to use these operations, only when needed. For example, let's say we've performed some kind of reduction on multiple blocks locally. We know that if we've used some shared of memory on this reduction, one must in addition add the block results, to get the full answer of our reduction. One way to do that is to force an atomic operation in all blocks, so they store it in the first address of the global memory, that we will later *consume* from the CPU. So if one suppose that we've accumulated the thread-local reduce into `sum` and the pointer to the initial data `g_out`, we would do something like

```
if(this_thread_block().thread_rank()==0){
    atomicAdd(&g_out[0], sum);
}
```

5.2.1 Asynchronous streams

We've already mentioned the asynchronous memory transfer mechanism, when discussing the memory models (see ??). Let's get back to the same example, and illustrate in more detail on what was happening in the codes.

```

1  for (int i = 0; i < nStreams; ++i){
2      int offset = i * streamSize; //offset for memory copy
3      cudaMemcpyAsync(&d_a[offset], &a[offset], streamBytes,\
4          cudaMemcpyHostToDevice, stream[i]);
5      kernel<<<streamSize/blockSize, blockSize, 0,\
6          stream[i]>>>(d_a, offset);
7      cudaMemcpyAsync(&a[offset], &d_a[offset], streamBytes,\
8          cudaMemcpyDeviceToHost, stream[i]);
9  }

```




Figure 9: Scenario of the execution pipeline, if no streams and asynchronous memory are involved.

Without the streams and the asynchronous memory copy, all the operations invoked in the loop would have been sequential. Suppose that one of requirements is not fulfilled. That is, suppose one replaces the `cudaMemcpyAsync()` with something else, and/or remove the stream functionality. Then the iteration of $i=1$ will wait for the $i=0$ iteration and so on (see ??).

One may see that the sequential version does not fully use the potential of the GPU. That is, some of the memory transfer mechanisms are waiting for others to finish. To solve this problem, one use the asynchronous memory version (??). In this case, the execution pipeline would look like that:

Note that this 10 pipeline is not always true. That is, based on different architectures, different results can be expected. For example, some of the versions have only a single memory copy and kernel engines. This implies that the async functionality does not make



Figure 10: Scenario of the execution pipeline, which is fully asynchronous and fully uses the memory transfer mechanism's resources.

sense for these versions (note that these versions are not very common anymore).

6 Appendix

Improving with primitive operations

The modulo operation `%` is expensive for an arithmetic operation. One could replace it by something much easier for the GPU architecture to execute. In this case, with a bit of knowledge of binary number representation, we can verify that `if((id_x%(stride*2)) == 0)` is equivalent to `if((id_x&(stride*2 -1)) == 0)`. It would impossible for me to come up with this small optimization on my own. Maybe for CS majors, it is evident.

References

- [1] Cuda c++ programming guide. URL: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#programming-model>.
- [2] Paraleljnoje programirovanije s cuda. chastj 1: Vvedenije, Apr 2015. URL: https://habr.com/ru/company/epam_systems/blog/245503/.
- [3] Cooperative groups: Flexible cuda thread programming, Aug 2020. URL: <https://developer.nvidia.com/blog/cooperative-groups/>.
- [4] Fang's notebook. nichijou.co, May 2022. URL: <https://nichijou.co/cudaRandom-memAlign/>.
- [5] Brijendar BakchodiaBrijendar Bakchodia1 and AmadanAmadan. What are atomic operations for newbies?, Apr 2018. URL: <https://stackoverflow.com/questions/52196678/what-are-atomic-operations-for-newbies>.
- [6] Mark Harris. How to overlap data transfers in cuda c/c++, Dec 2013. URL: <https://developer.nvidia.com/blog/how-overlap-data-transfers-cuda-cc/>.
- [7] Raj Prasanna Ponnuraj. Cuda-memory model, Oct 2020. URL: <https://medium.com/analytics-vidhya/cuda-memory-model-823f02cef0bf#:~:text=Pageable%20memory,-The%20memory%20allocated&text=The%20data%20at%20this%20memory,allocation%20and%20transfer%20is%20slow>.
- [8] Brian Tuomanen. *Hands-On GPU Programming with Python and CUDA: Explore high-performance parallel computing with CUDA*. Packt Publishing Ltd, 2018.